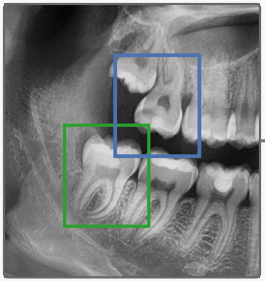
(a) Grounding taken as an answer



returned box (blue), expert box (green)



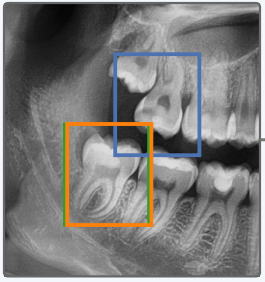
crop of the returned box

The crop looks authoritative.

Every later observation is built on a box that is wrong - here it sits 0.92 box-sizes from the expert box (IoU 0.07).

Nothing in the loop questions the box.

(b) Grounding as a hypothesis (ours)



same episode; box corrected (orange)

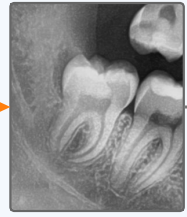


inspect: “... about one arch too high”

nudge_crop × 1

dx –0.57 · dy +0.69

IoU 0.07 → 0.94



re-inspect: after nudge_crop

final answer

Q4T8 · Impacted tooth
confidence 0.92

matches expert label

box returned by the grounding tool

expert box

box after nudge_crop