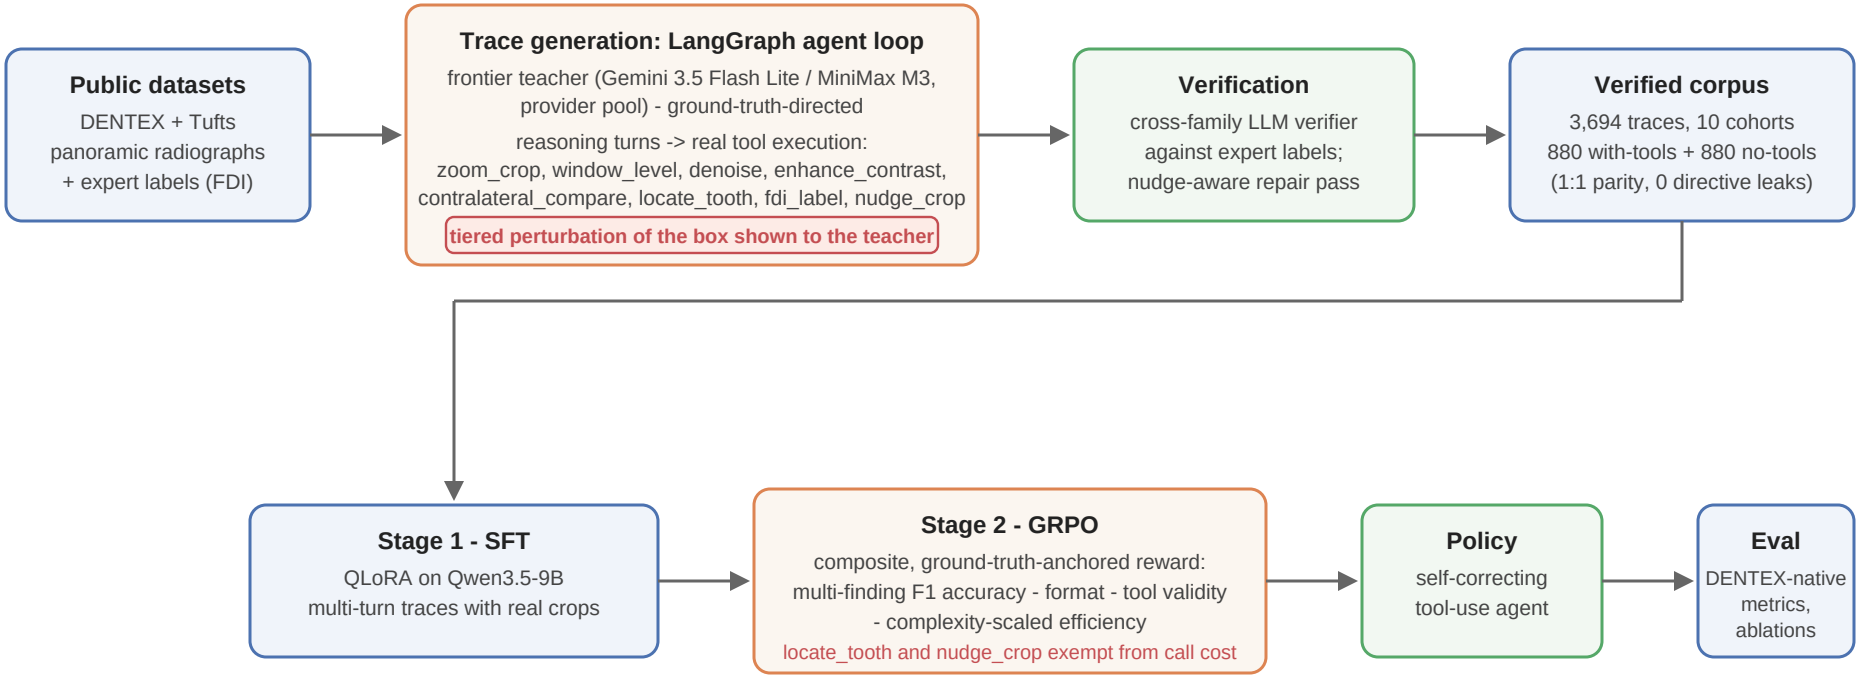

Both training and evaluation run the same agent loop and the same image-input geometry (native, or canonical views when training on TPU); tool calls execute for real against the source image at every stage.
During trace generation only, locate_tooth returns the expert-verified box for annotated teeth, and the perturbation hook independently perturbs what the teacher is shown.